

Chapter 22

Transaction Management (con't)

Granularity of Data Items

- ◆ **Size of data items chosen as unit of protection by concurrency control protocol.**
- ◆ **Ranging from coarse to fine:**
 - **The entire database.**
 - **A file.**
 - **A page (or area or database spaced).**
 - **A record.**
 - **A field value of a record.**

Granularity of Data Items

- ◆ **Tradeoff:**
 - coarser, the lower the degree of concurrency;
 - finer, more locking information that is needed to be stored.
- ◆ **Best item size depends on the types of transactions.**

Hierarchy of Granularity

- ◆ Could represent granularity of locks in a hierarchical structure.
- ◆ Root node represents entire database, level 1s represent files, etc.
- ◆ When node is locked, all its descendants are also locked.
- ◆ DBMS should check hierarchical path before granting lock.

Hierarchy of Granularity

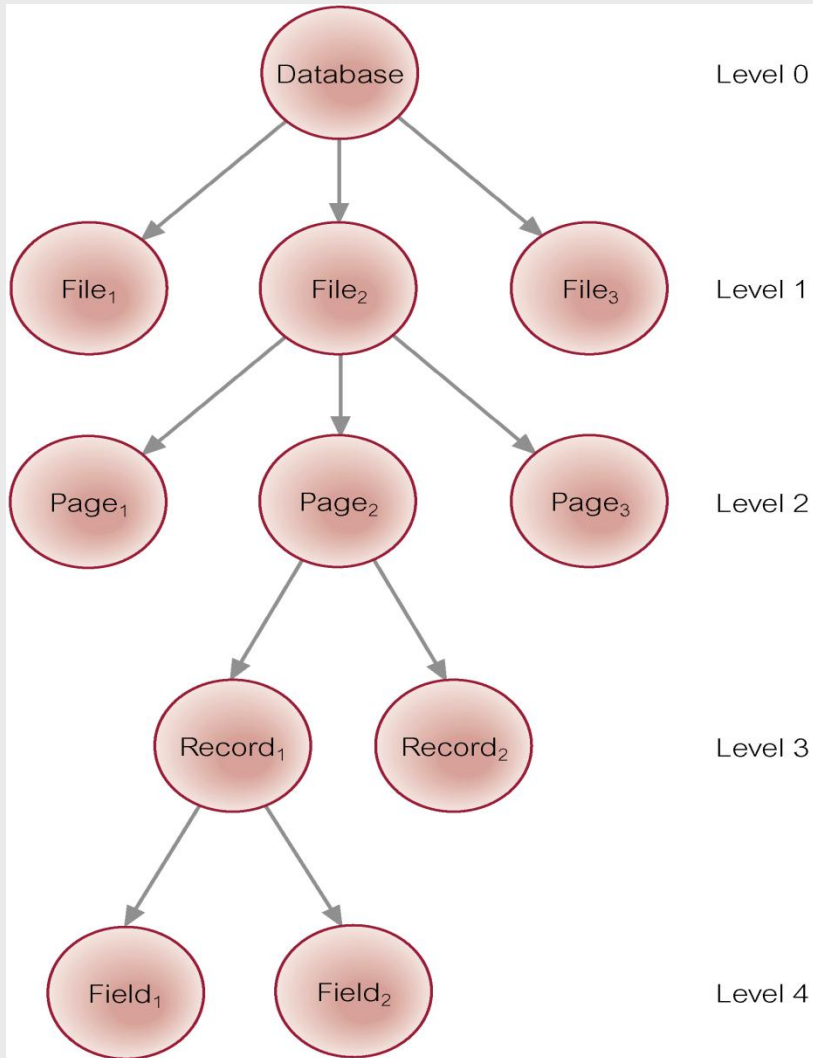
- ◆ *Intention lock* could be used to lock all ancestors of a locked node.
- ◆ Intention locks can be read or write. Applied top-down, released bottom-up.

Table 20.1 Lock compatibility table for multiple-granularity locking.

	IS	IX	S	SIX	X
<i>IS</i>	✓	✓	✓	✓	✗
<i>IX</i>	✓	✓	✗	✗	✗
<i>S</i>	✓	✗	✓	✗	✗
<i>SIX</i>	✓	✗	✗	✗	✗
<i>X</i>	✗	✗	✗	✗	✗

✓ = compatible, ✗ = incompatible

Levels of Locking



Database Recovery

Process of restoring database to a correct state in the event of a failure.

◆ Need for Recovery Control

- Two types of storage: volatile (main memory) and nonvolatile.**
- Volatile storage does not survive system crashes.**
- Stable storage represents information that has been replicated in several nonvolatile storage media with independent failure modes.**

Types of Failures

- ◆ **System crashes, resulting in loss of main memory.**
- ◆ **Media failures, resulting in loss of parts of secondary storage.**
- ◆ **Application software errors.**
- ◆ **Natural physical disasters.**
- ◆ **Carelessness or unintentional destruction of data or facilities.**
- ◆ **Sabotage.**

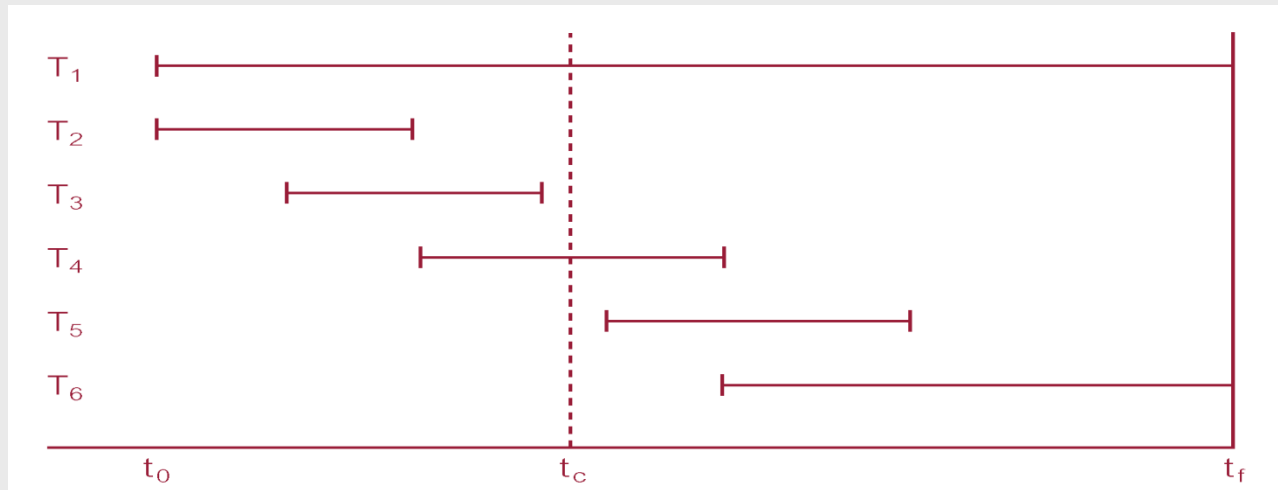
Transactions and Recovery

- ◆ Transactions represent basic unit of recovery.
- ◆ Recovery manager responsible for atomicity and durability.
- ◆ If failure occurs between commit and database buffers being flushed to secondary storage then, to ensure durability, recovery manager has to *redo (rollforward)* transaction's updates.

Transactions and Recovery

- ◆ If transaction had not committed at failure time, recovery manager has to *undo (rollback)* any effects of that transaction for atomicity.
- ◆ Partial undo - only one transaction has to be undone.
- ◆ Global undo - all transactions have to be undone.

Example



- ◆ DBMS starts at time t_0 , but fails at time t_f . Assume data for transactions T_2 and T_3 have been written to secondary storage.
- ◆ T_1 and T_6 have to be undone. In absence of any other information, recovery manager has to redo T_2 , T_3 , T_4 , and T_5 .

Recovery Facilities

- ◆ **DBMS should provide following facilities to assist with recovery:**
 - **Backup mechanism, which makes periodic backup copies of database.**
 - **Logging facilities, which keep track of current state of transactions and database changes.**
 - **Checkpoint facility, which enables updates to database in progress to be made permanent.**
 - **Recovery manager, which allows DBMS to restore database to consistent state following a failure.**

Log File

- ◆ Contains information about all updates to database:
 - Transaction records.
 - Checkpoint records.
- ◆ Often used for other purposes (for example, auditing).

Log File

- ◆ **Transaction records contain:**
 - **Transaction identifier.**
 - **Type of log record, (transaction start, insert, update, delete, abort, commit).**
 - **Identifier of data item affected by database action (insert, delete, and update operations).**
 - **Before-image of data item.**
 - **After-image of data item.**
 - **Log management information.**

Sample Log File (This log file not related to Slide 11)

Tid	Time	Operation	Object	Before image	After image	pPtr	nPtr
T1	10:12	START				0	2
T1	10:13	UPDATE	STAFF SL21	(old value)	(new value)	1	8
T2	10:14	START				0	4
T2	10:16	INSERT	STAFF SG37		(new value)	3	5
T2	10:17	DELETE	STAFF SA9	(old value)		4	6
T2	10:17	UPDATE	PROPERTY PG16	(old value)	(new value)	5	9
T3	10:18	START				0	11
T1	10:18	COMMIT				2	0
	10:19	CHECKPOINT	T2, T3				
T2	10:19	COMMIT				6	0
T3	10:20	INSERT	PROPERTY PG4		(new value)	7	12
T3	10:21	COMMIT				11	0

Log File

- ◆ Log file may be duplexed or triplexed.
- ◆ Log file sometimes split into two separate random-access files.
- ◆ Potential bottleneck; critical in determining overall performance.

Checkpointing

Checkpoint

Point of synchronization between database and log file. All buffers are force-written to secondary storage.

- ◆ **Checkpoint record is created containing identifiers of all active transactions.**
- ◆ **When failure occurs, redo all transactions that committed since the checkpoint and undo all transactions active at time of crash.**

Checkpointing

- ◆ In previous example, with checkpoint at time t_c , changes made by T_2 and T_3 have been written to secondary storage.
- ◆ Thus:
 - only redo T_4 and T_5 ,
 - undo transactions T_1 and T_6 .

Recovery Techniques

- ◆ **If database has been damaged:**
 - Need to restore last backup copy of database and reapply updates of committed transactions using log file.
- ◆ **If database is only inconsistent:**
 - Need to undo changes that caused inconsistency. May also need to redo some transactions to ensure updates reach secondary storage.
 - Do not need backup, but can restore database using before- and after-images in the log file.

Main Recovery Techniques

- ◆ **Three main recovery techniques:**
 - **Deferred Update**
 - **Immediate Update**
 - **Shadow Paging**

Deferred Update

- ◆ Updates are not written to the database until after a transaction has reached its commit point.
- ◆ If transaction fails before commit, it will not have modified database and so no undoing of changes required.
- ◆ May be necessary to redo updates of committed transactions as their effect may not have reached database.

Immediate Update

- ◆ Updates are applied to database as they occur.
- ◆ Need to redo updates of committed transactions following a failure.
- ◆ May need to undo effects of transactions that had not committed at time of failure.
- ◆ Essential that log records are written before write to database. *Write-ahead log protocol.*

Immediate Update

- ◆ If no “*transaction commit*” record in log, then that transaction was active at failure and must be undone.
- ◆ Undo operations are performed *in reverse order in which they were written to log*.

Shadow Paging

- ◆ Maintain two page tables during life of a transaction: *current* page and *shadow* page table.
- ◆ When transaction starts, two pages are the same.
- ◆ Shadow page table is never changed thereafter and is used to restore database in event of failure.
- ◆ During transaction, current page table records all updates to database.
- ◆ When transaction completes, current page table becomes shadow page table.