

Chapter 24

Distributed DBMSs – Concepts and Design

Chapter 24 - Objectives

- ◆ Concepts.
- ◆ Advantages and disadvantages of distributed databases.
- ◆ Functions and architecture for a DDBMS.
- ◆ Distributed database design.
- ◆ Levels of transparency.
- ◆ Comparison criteria for DDBMSs.

Concepts

Distributed Database

A logically interrelated collection of shared data (and a description of this data), physically distributed over a computer network.


Distributed DBMS

Software system that permits the management of the distributed database and makes the distribution transparent to users.

Concepts

- ◆ **Collection of logically-related shared data.**
- ◆ **Data split into fragments.**
- ◆ **Fragments may be replicated.**
- ◆ **Fragments/replicas allocated to sites.**
- ◆ **Sites linked by a communications network.**
- ◆ **Data at each site is under control of a DBMS.**
- ◆ **DBMSs handle local applications autonomously.**
- ◆ **Each DBMS participates in at least one global application.**

Distributed DBMS



Types of DDBMS

- ◆ **Homogeneous DDBMS**
- ◆ **Heterogeneous DDBMS**

Homogeneous DDBMS

- ◆ All sites use same DBMS product.
- ◆ Much easier to design and manage.
- ◆ Approach provides incremental growth and allows increased performance.

Heterogeneous DDBMS

- ◆ Sites may run different DBMS products, with possibly different underlying data models.
- ◆ Occurs when sites have implemented their own databases and integration is considered later.
- ◆ Translations required to allow for:
 - Different hardware.
 - Different DBMS products.
 - Different hardware and different DBMS products.
- ◆ Typical solution is to use *gateways*.

Functions of a DDBMS

- ◆ **Expect DDBMS to have at least the functionality of a DBMS.**
- ◆ **Also to have following functionality:**
 - **Extended communication services.**
 - **Extended Data Dictionary.**
 - **Distributed query processing.**
 - **Extended concurrency control.**
 - **Extended recovery services.**

Distributed Database Design

- ◆ **Three key issues:**
 - **Fragmentation,**
 - **Allocation,**
 - **Replication.**

Distributed Database Design

Fragmentation

Relation may be divided into a number of sub-relations, which are then distributed.

Allocation

Each fragment is stored at site with “optimal” distribution.

Replication

Copy of fragment may be maintained at several sites.

Fragmentation

- ◆ **Definition and allocation of fragments carried out strategically to achieve:**
 - **Locality of Reference.**
 - **Improved Reliability and Availability.**
 - **Improved Performance.**
 - **Balanced Storage Capacities and Costs.**
 - **Minimal Communication Costs.**
- ◆ **Involves analyzing most important applications, based on quantitative/qualitative information.**

Fragmentation

- ◆ Quantitative information may include:
 - frequency with which an application is run;
 - site from which an application is run;
 - performance criteria for transactions and applications.
- ◆ Qualitative information may include transactions that are executed by application, type of access (read or write), and predicates of read operations.

Data Allocation

- ◆ **Four alternative strategies regarding placement of data:**
 - **Centralized,**
 - **Partitioned (or Fragmented),**
 - **Complete Replication,**
 - **Selective Replication.**

Data Allocation

Centralized: Consists of single database and DBMS stored at one site with users distributed across the network.

Partitioned: Database partitioned into disjoint fragments, each fragment assigned to one site.

Complete Replication: Consists of maintaining complete copy of database at each site.

Selective Replication: Combination of partitioning, replication, and centralization.

Why Fragment?

◆ Usage

- Applications work with views rather than entire relations.

◆ Efficiency

- Data is stored close to where it is most frequently used.
- Data that is not needed by local applications is not stored.

Why Fragment?

◆ Parallelism

- With fragments as unit of distribution, transaction can be divided into several subqueries that operate on fragments.

◆ Security

- Data not required by local applications is not stored and so not available to unauthorized users.

Why Fragment?


◆ Disadvantages

- Performance,
- Integrity.



Horizontal and Vertical Fragmentation



(a)



(b)



Correctness of Fragmentation

◆ Three correctness rules:

- Completeness,
- Reconstruction,
- Disjointness.

Correctness of Fragmentation

Completeness

If relation R is decomposed into fragments R_1, R_2, \dots, R_n , each data item that can be found in R must appear in at least one fragment.

Reconstruction

- ◆ Must be possible to define a relational operation that will reconstruct R from the fragments.
- ◆ Reconstruction for horizontal fragmentation is Union operation and Join for vertical .

Correctness of Fragmentation

Disjointness

- ◆ If data item d_i appears in fragment R_i , then it should not appear in any other fragment.
- ◆ Exception: vertical fragmentation, where primary key attributes must be repeated to allow reconstruction.
- ◆ For horizontal fragmentation, data item is a tuple.
- ◆ For vertical fragmentation, data item is an attribute.

Types of Fragmentation

- ◆ **Four types of fragmentation:**

- **Horizontal,**
- **Vertical,**
- **Mixed,**
- **Derived.**


- ◆ **Other possibility is no fragmentation:**

- **If relation is small and not updated frequently, may be better not to fragment relation.**



Horizontal and Vertical Fragmentation




(a)




(b)



Mixed Fragmentation



(a)



(b)

Horizontal Fragmentation

- ◆ Consists of a subset of the tuples of a relation.
- ◆ Defined using *Selection* operation of relational algebra:

$$\sigma_p(R)$$

- ◆ For example:

$$P_1 = \sigma_{\text{type}='House'}(\text{PropertyForRent})$$

$$P_2 = \sigma_{\text{type}='Flat'}(\text{PropertyForRent})$$

Vertical Fragmentation

- ◆ Consists of a subset of attributes of a relation.
- ◆ Defined using *Projection* operation of relational algebra:

$$\Pi_{a_1, \dots, a_n}(R)$$

- ◆ For example:

$$S_1 = \Pi_{\text{staffNo}, \text{position}, \text{sex}, \text{DOB}, \text{salary}}(\text{Staff})$$

$$S_2 = \Pi_{\text{staffNo}, \text{fName}, \text{lName}, \text{branchNo}}(\text{Staff})$$

- ◆ Determined by establishing *affinity* of one attribute to another.

Distributed Database Design Methodology

- 1. Use normal methodology to produce a design for the global relations.**
- 2. Examine topology of system to determine where databases will be located.**
- 3. Analyze most important transactions and identify appropriateness of horizontal/vertical fragmentation.**
- 4. Decide which relations are not to be fragmented.**

Transparencies in a DDBMS

◆ Distribution Transparency

- Fragmentation Transparency
- Location Transparency
- Replication Transparency
- Local Mapping Transparency

Transparencies in a DDBMS

- ◆ **Transaction Transparency**
 - **Concurrency Transparency**
 - **Failure Transparency**
- ◆ **DBMS Transparency**

Distribution Transparency

- ◆ **Distribution transparency allows user to perceive database as single, logical entity.**
- ◆ **If DDBMS exhibits distribution transparency, user does not need to know:**
 - data is fragmented (fragmentation transparency),
 - location of data items (location transparency),
 - otherwise call this local mapping transparency.
- ◆ **With replication transparency, user is unaware of replication of fragments .**

Naming Transparency

- ◆ Each item in a DDB must have a unique name.
- ◆ DDBMS must ensure that no two sites create a database object with same name.
- ◆ One solution is to create central name server. However, this results in:
 - loss of some local autonomy;
 - central site may become a bottleneck;
 - low availability; if the central site fails, remaining sites cannot create any new objects.

Transaction Transparency

- ◆ Ensures that all distributed transactions maintain distributed database's integrity and consistency.
- ◆ Distributed transaction accesses data stored at more than one location.
- ◆ Each transaction is divided into number of subtransactions, one for each site that has to be accessed.
- ◆ DDBMS must ensure the indivisibility of both the global transaction and each of the subtransactions.

Concurrency Transparency

- ◆ All transactions must execute independently and be logically consistent with results obtained if transactions executed one at a time, in some arbitrary serial order.
- ◆ Same fundamental principles as for centralized DBMS.
- ◆ DDBMS must ensure both global and local transactions do not interfere with each other.
- ◆ Similarly, DDBMS must ensure consistency of all subtransactions of global transaction.

Concurrency Transparency

- ◆ Results of concurrent transactions same as some serial order of transaction execution.
- ◆ Replication makes concurrency more complex.
- ◆ If a copy of a replicated data item is updated, update must be propagated to all copies.
- ◆ Could propagate changes as part of original transaction, making it an atomic operation.
- ◆ However, if one site holding copy is not reachable, then transaction is delayed until site is reachable.

Concurrency Transparency

- ◆ Could limit update propagation to only those sites currently available. Remaining sites updated when they become available again.
- ◆ Could allow updates to copies to happen asynchronously, sometime after the original update. Delay in regaining consistency may range from a few seconds to several hours.

Failure Transparency

- ◆ **DDBMS must ensure atomicity and durability of global transaction.**
- ◆ **Means ensuring that subtransactions of global transaction either all commit or all abort.**
- ◆ **Thus, DDBMS must synchronize global transaction to ensure that all subtransactions have completed successfully before recording a final COMMIT for global transaction.**
- ◆ **Must do this in presence of site and network failures.**